

PWN College

Session 13

Atousa Ahsani

References: <https://pwn.college/>, <https://guyinatuxedo.github.io/>

Stack Buffer Overflows

Stack Canary

Relro

Stack Canary

- The **Stack Canary** is another mitigation designed to protect against things like **stack based** buffer overflows.
- The general idea is, a **random value** is placed at the bottom of the stack frame, which is **below** the **stack variables** where we actually have input.
- If had a buffer overflow to overwrite the saved return address, this value on the stack would be **overwritten**. Then before the return address is executed, it checks to see if that value is the same one it set. If it isn't then it knows that there is a **memory corruption** bug happening and terminates the program.
- Also the name comes from the use of **canaries** in a **mine**. If the canary stops singing, get out before you die from gas poisoning.

Stack Canary

- During compilation, the compiler will insert a **canary** check stub. The following options are supported by any recent compiler:
 - *-fstack-protector* (since GCC 4.1): includes a **canary** when a function defines an **array of char** with a size of **8 bytes or more**
 - *-fstack-protector-all*: adds a canary for all non-inline functions
 - *-fstack-protector-strong* (since 4.9): provides a smarter way to protect **any sensitive location** within the current context.

Stack Canary

- let's look at a binary compiled with a stack canary.

```
→ canary gcc -fstack-protector-all main.c -o canary0n
→ canary gdb canary0n
Reading symbols from canary0n...
(No debugging symbols found in canary0n)
gdb-peda$ disas main
Dump of assembler code for function main:
0x0000000000001169 <+0>:    endbr64
0x000000000000116d <+4>:    push    rbp
0x000000000000116e <+5>:    mov     rbp, rsp
0x0000000000001171 <+8>:    sub     rsp, 0x10
0x0000000000001175 <+12>:   mov     rax, QWORD PTR fs:0x28
0x000000000000117e <+21>:   mov     QWORD PTR [rbp-0x8], rax
0x0000000000001182 <+25>:   xor     eax, eax
0x0000000000001184 <+27>:   lea     rax, [rbp-0xc]
0x0000000000001188 <+31>:   mov     rsi, rax
0x000000000000118b <+34>:   lea     rdi, [rip+0xe72]          # 0x2004
0x0000000000001192 <+41>:   mov     eax, 0x0
0x0000000000001197 <+46>:   call    0x1070 <__isoc99_scanf@plt>
0x000000000000119c <+51>:   mov     eax, 0x0
0x00000000000011a1 <+56>:   mov     rdx, QWORD PTR [rbp-0x8]
0x00000000000011a5 <+60>:   xor     rdx, QWORD PTR fs:0x28
0x00000000000011ae <+69>:   je      0x11b5 <main+76>
0x00000000000011b0 <+71>:   call    0x1060 <stack_chk_fail@plt>
0x00000000000011b5 <+76>:   leave
0x00000000000011b6 <+77>:   ret
End of assembler dump.
```

Stack Canary

- Now let's look at a binary compiled from the same source code, but **without** a stack canary.

```
→ canary gcc -fno-stack-protector main.c -o canaryOff
→ canary gdb canaryOff
Reading symbols from canaryOff...
(No debugging symbols found in canaryOff)
gdb-peda$ disas main
Dump of assembler code for function main:
0x0000000000001149 <+0>:      endbr64
0x000000000000114d <+4>:      push    rbp
0x000000000000114e <+5>:      mov     rbp, rsp
0x0000000000001151 <+8>:      sub     rsp, 0x10
0x0000000000001155 <+12>:     lea     rax, [rbp-0x4]
0x0000000000001159 <+16>:     mov     rsi, rax
0x000000000000115c <+19>:     lea     rdi, [rip+0xea1]          # 0x2004
0x0000000000001163 <+26>:     mov     eax, 0x0
0x0000000000001168 <+31>:     call   0x1050 <__isoc99_scanf@plt>
0x000000000000116d <+36>:     mov     eax, 0x0
0x0000000000001172 <+41>:     leave
0x0000000000001173 <+42>:     ret
End of assembler dump.
```

- We can see a few **differences** between the code, like when it **checks** the **stack canary**.

Stack Canary

- **How can we recognize the canary?**
- We can tell that a value is the **stack canary** from several different things.
 - Firstly it is the **value** being used when it is doing the **stack canary check**.
 - Also it is around the spot on the stack it should be.
 - Also it matches the **pattern** of a **stack canary**. While they are random they do fit a general pattern.
- **What is the pattern?**
 - For **x64 elfs**, the pattern is an **0x8 byte qword**, where the **first seven bytes** are **random** and the **last byte** is a **null byte**.
 - For **x86 elfs**, the pattern is a **0x4 byte dword**, where the **first three bytes** are **random** and the **last byte** is a **null byte**.

Stack Canary

- let's examine a **64-bit** binary. We set **two breakpoints**. One exactly **before** the **scanning input** and the other one right **after** that.
- Program will be **interrupted** before scanning in.
- Here we just wrote the value of the **canary** to itself, and it **passed** the check. Of course this requires us to know the value of the stack canary. This can be accomplished via **leaking** the **canary** (which we will see later).
- Also in some cases you might be able to do something like **brute forcing** that value.

```
→ canary gdb canary0n
gdb-peda$ b* main+46
Breakpoint 1 at 0x1197
gdb-peda$ b* main+51
Breakpoint 2 at 0x119c
```

```
gdb-peda$ r
.
.
.
Breakpoint 1, 0x0000555555551197 in main ()
gdb-peda$ x/8w $rsp
0x7fffffffdded0: 0xfffffffffd0  0x000007fff  0x3c083100  0x7f4e6530
0x7fffffffdee0: 0x000000000  0x000000000  0xfdde0b3  0x000007fff
gdb-peda$ c
Continuing.
aaaaaaaaaaaaaaaaaaaa => This input causes stack smashing!
.
.
Breakpoint 2, 0x000055555555119c in main ()
gdb-peda$ x/8w $rsp
0x7fffffffdded0: 0xfffffffffd0  0x61616161  0x61616161  0x61616161
0x7fffffffdee0: 0x61616161  0x61616161  0x00616161  0x000007fff
gdb-peda$ set *0x7fffffffdded8 = 0x3c083100
gdb-peda$ set *0x7fffffffddedc = 0x7f4e6530
gdb-peda$ x/8w $rsp
0x7fffffffdded0: 0xfffffffffd0  0x61616161  0x3c083100  0x7f4e6530
0x7fffffffdee0: 0x61616161  0x61616161  0x00616161  0x000007fff
gdb-peda$ c
Continuing.
.
.
Invalid $PC address: 0x7fff00616161
There is no stack smashing error!
We overwrote the canary!
```


Stack Buffer Overflows

Stack Canary

Relro

Relro

- **Relro (Read only Relocation)** affects the **memory permissions** similar to **NX**.
- The difference is whereas with **NX** it makes the stack **executable**, **RELRO** makes certain things **read only** so we can't write to them.
- The most common way I've seen this be an obstacle is **preventing us** from doing a ***got*** table overwrite, which will be covered later.

Relro

- A **dynamically linked ELF** binary uses a **look-up table** called the **Global Offset Table (GOT)** to dynamically resolve functions that are located in **shared libraries**.
- Such **calls** point to the **Procedure Linkage Table (PLT)**, which is present in the *.plt* section of the binary. The *.plt* section contains **x86 instructions** that point directly to the **GOT**, which lives in the *.got.plt* section.
- **GOT** normally contains pointers that point to the **actual location** of these functions in the **shared libraries** in **memory**.
- The **GOT** is populated **dynamically** as the program is running.
- The first time a **shared function** is called, program jumps to **PTL** section. Then the code is followed by a jump to **GOT** section. Since the **GOT** has no addresses for the first time, the **dynamic linker** is called to find the actual location of the function in question. The location found is then written to the **GOT**.
- The second time a function is called, the **GOT** contains the known location of the function. This is called "**lazy binding**." This is because it is unlikely that the location of the shared function has changed and it saves some CPU cycles as well.

Relro

- There are a few implications of the above.
- Firstly, **PLT** needs to be located at a fixed offset from the **.text** section.
- Secondly, since **GOT** contains data used by different parts of the program directly, it needs to be allocated at a known **static address** in **memory**.
- Lastly, and more importantly, because the **GOT** is lazily bound it needs to be writable.
- Since **GOT** exists at a **predefined place** in memory, a program that contains a vulnerability allowing an attacker to write 4 bytes at a controlled place in memory (such as some integer overflows leading to out-of-bounds write), may be exploited to allow arbitrary code execution.

Relro

- To prevent the above mentioned security weakness, we need to ensure that the **linker** resolves **all dynamically linked functions** at the **beginning** of the execution, and then makes the **GOT read-only**.
- This technique is called **RELRO** and ensures that the **GOT** cannot be overwritten in vulnerable ELF binaries.
- RELRO can be turned on when compiling a program by using the following options:
 - Full RELRO:

```
gcc -g -Wl,-z,relro,-z,now -o test testcase.c
```

- Partial RELRO:

```
gcc -g -Wl,-z,relro -o test testcase.c
```

Relro

- In **partial RELRO**, the **non-PLT** part of the **GOT** section (.got from readelf output) is **read only** but **.got.plt** is still **writable**.
- Whereas in full RELRO, the entire **GOT** (.got and .got.plt both) is marked as **read-only**.
- Let's see what the **memory permissions** look like for a **got** table entry for a binary **with** and **without** relro.
- We will use this simple code.

```
int main(int argc, char const *argv[])
{
    printf("Hello everyone!\n");
    printf("This is an example!\n");
    exit(0);
    return 0;
}
```

Relro

- With relro

- Compilation: `gcc -g -Wl,-z,relro,-z,now -o full main.c`
- Gef gdb:

```
gef> p puts
$1 = {int (const char *)} 0x7ffff7e3e5a0 <__GI_IO_puts>
gef> search-pattern 0x7ffff7e3e5a0
[+] Searching '\xa0\xe5\xe3\xf7\xff\xf7' in memory
[+] In '/home/atousa/PWNCollegeCourse_TMU/13/relro/full' (0x555555557000-0x555555558000), permission=r--
    0x555555557fc8 - 0x555555557fe0 -> "\xa0\xe5\xe3\xf7\xff\xf7[...]"
```

- Without relro:

- Compilation: `gcc -g -Wl,-z,norelro -o norelro main.c`
- Gef gdb:

```
gef> p puts
$1 = {int (const char *)} 0x7ffff7e3e5a0 <__GI_IO_puts>
gef> search-pattern 0x7ffff7e3e5a0
[+] Searching '\xa0\xe5\xe3\xf7\xff\xf7' in memory
[+] In '/home/atousa/PWNCollegeCourse_TMU/13/relro/norelro' (0x555555557000-0x555555558000), permission=rw-
    0x555555557390 - 0x5555555573a8 -> "\xa0\xe5\xe3\xf7\xff\xf7[...]"
```